

SUMMARY

Experienced Software Developer with 4+ years of success building scalable backend systems and cloud-native applications. Skilled in Java, Python, Flask, React, and PostgreSQL, with strong experience in AWS and Google Cloud. Specialized in developing GenAI-powered tools, integrating LLMs into production services, and optimizing performance. Adept at creating secure, maintainable software and collaborating in Agile teams to deliver reliable, real-world solutions across healthcare and enterprise domains.

SKILLS

- **Programming Languages:** Java, Python, JavaScript, SQL, C++, Dart, Kotlin, HTML5, CSS
- **Backend & Cloud:** Spring Boot, FastAPI, Flask, Node.js, Express.js, PEGA, RESTful APIs, JWT Authentication, Firebase (Auth, Firestore, Realtime DB, Messaging, Cloud Storage), AWS (Amplify, Lambda, API Gateway, Cognito, S3, DynamoDB, SNS, CloudWatch), Google Cloud (Firestore, Functions, App Distribution)
- **Frontend & Mobile Development:** React.js, Flutter, Android Studio, Jetpack Components, XML UI, Room Database, Component-Based Architecture, Responsive Design, Provider/Bloc, MVVM Architecture, Material Design
- **Databases:** PostgreSQL, MySQL, SQL Server, MongoDB, DynamoDB
- **DevOps & Monitoring:** Docker, Kubernetes, Terraform, CI/CD Pipelines (GitHub Actions, Gradle), Git, JIRA, Splunk, AppDynamics, Grafana, Crash Reporting
- **AI & Machine Learning:** Generative AI, Large Language Models (LLMs), Prompt Engineering, OpenAI APIs, LangChain, Hugging Face Transformers, TensorFlow, TensorFlow Lite, Scikit-learn, NLP (NER, Text Summarization, Sentiment Analysis), Recommendation Systems, Model Deployment, Data Analysis & Visualization
- **Tools & Platforms:** IntelliJ, VS Code, Postman, Figma, Wireframing, Prototyping, GitHub Copilot, Cursor
- **Methodologies & Core Competencies:** Agile/Scrum, SDLC, Design Patterns, HIPAA Compliance, OOP, Debugging, Performance Optimization, Problem Solving, Team Collaboration, Time Management

PROFESSIONAL EXPERIENCE

Software Developer | Tempus

Oct 2024 – Present | USA

- Developed secure, high-throughput backend services using **FastAPI** and integrated them with **LLMs via OpenAI GPT-4 and LangChain**, enabling oncology researchers to auto-summarize 2TB+ of unstructured patient notes in under 10 seconds per request.
- Built a **GenAI-enabled clinical documentation assistant**, deployed on **GCP using Kubernetes and Docker**, which reduced manual reporting time by ~4 hours per oncologist weekly through prompt-engineered summarization and keyword extraction.
- Designed and implemented **Java-based orchestration logic** to interface with AI pipelines, providing queue-based job handling for LLM inference tasks and ensuring consistent delivery under high concurrency.
- Collaborated with data engineering teams to create **tokenized RESTful APIs in Java 17 and Spring Boot**, enforcing **JWT authentication**, rate limiting, and input sanitization for HIPAA-compliant data access.
- Led the refactoring and modularization of existing Java microservices, improving unit test coverage from 62% to 91%, reducing code complexity, and cutting deployment-related bugs by 45% over two release cycles.

Software Development Associate – Client: T-Mobile | Deloitte

May 2020 – Aug 2023 | India

- Designed and deployed 40+ **RESTful** microservices using Java 8, Spring Boot, and Spring Security, integrated with Oracle/**PostgreSQL**, processing 120K+ daily transactions across telecom systems.
- Engineered backend modules for asynchronous processing with **JMS** and **Java Mail API**, enabling real-time alerts and batch notifications, improving incident response by 38%.
- Developed internal web applications with **Java-based APIs**, ReactJS, **HTML5**, and Bootstrap, reducing ticket resolution time by 25% for support teams.
- Automated cloud operations with custom Java agents, streamlining **AWS S3** cleanups and system health checks, cutting manual intervention by 60+ hours/month.
- Optimized multi-table joins and transactional queries via **Spring Data JPA** custom finder methods, reducing database response times by 30% on high-traffic endpoints.

EDUCATION

Florida Institute of Technology

Aug 2023 - May 2025

Master's, Computer Engineering

CERTIFICATIONS

- AWS Certified Developer – Associate
- Introduction to Large Language Models
- Supervised Machine Learning: Regression
- Introduction to Generative AI